



Product Case Study

ArgusAI

From alert to root cause, in one screen.

AI-powered incident investigation for software teams

Irfan Ahmed J
Product Strategy & Product Design



INCIDENT FEED

ALERT · Critical

High latency on checkout-service

Fired 14:02 · Alertmanager



AI ROOT-CAUSE ANALYSIS

Probable cause: database connection pool exhaustion


Evidence: metrics + logs

Actions: recommended next steps

Confidence: shown per analysis

Illustrative mockup. Replace with a real dashboard screenshot.

An alert is the start of the investigation, not the answer.



The person under pressure at 2 a.m. is the user; the team lead is the buyer.



Primary

On-call Engineer

JOB TO BE DONE

Restore service fast

FRUSTRATION

Unfamiliar service, scattered data



SECONDARY

Engineering manager / SRE lead

JOB TO BE DONE

Lower incident impact, fewer escalations

FRUSTRATION

Senior engineers become the bottleneck



TERTIARY

Junior / new engineer


JOB TO BE DONE

Contribute confidently to on-call

FRUSTRATION

Doesn't know where to look

Today's investigation is a manual loop of context-switching.



On-call engineers need to reach a **credible root-cause hypothesis** quickly, but the evidence is **scattered across tools** and interpreting it depends on **tribal knowledge**.

HOW MIGHT WE...

assemble the evidence and a first hypothesis before the engineer even opens a laptop?

The signals exist. The synthesis doesn't.

Data collection

Metrics and logs are gathered and displayed

Alerting

A threshold is crossed and someone is paged

Explanation

Correlate evidence across sources and say why → ArgusAI

Human decision

Engineer judges, acts, verifies

Why now

“Recent LLMs can reason over structured evidence, which makes automated first-pass analysis feasible.”

ArgusAI: an AI first-responder for incidents.



Detects

incidents from your existing alerts


Gathers

metrics, logs and context automatically

Explains

the probable cause, with evidence and next steps

From alert to action in one screen.



Steps 1–4: no human effort

Step 5: the engineer starts from a hypothesis, not a blank screen

Every feature exists to build speed and trust.

Speed

1

Automatic incident creation

Alerts become incidents with no manual step

2

Real-time push

WebSockets update the dashboard live

3

Auto-gathered evidence

Metrics and logs pulled for you

Trust

4

Evidence shown with the RCA

Engineers can verify, not just believe

5

Confidence score

Every analysis states how sure it is

6

Actions recommended, not automated


A human stays in control

Evidence in, explanation out.




Design choice: the model only sees retrieved evidence, to reduce unsupported claims.

Built on the tools teams already use.



On each alert, the AI service pulls the relevant metrics and logs, combines them with incident context, and asks Gemini for a structured RCA.


Problem → evidence → RCA → action.



Every choice traded something off.

Decision	Chose	Gave up	Why
Trigger	Automatic on alert	On-demand control, lower LLM cost	Speed matters most at incident time
Output	Recommend actions	Auto-remediation	Keep a human in the loop while trust is low
Transparency	Evidence + confidence shown	A simpler UI	Engineers won't act on unexplained AI output
Data scope	Metrics + logs only	Traces, deploys, change history	Ship a focused v1
Delivery	Real-time push	Simplicity of polling	Incidents are time-sensitive

Success is faster, trusted investigation, not just fast output.



Knowing where the product can fail is part of the product.



Next steps: breadth of evidence and depth of trust.

NOW

Feedback buttons on RCAs

Solves: Measures quality

Evidence redaction

Solves: Data sensitivity

User interviews

Solves: Validates the problem

NEXT

Deploy and change correlation

Solves: Limited context

Distributed traces

Solves: Limited context

Group related alerts

Solves: Alert storms

LATER

Similar past incidents

Solves: Faster diagnosis

Slack / PagerDuty integration

Solves: Fits the workflow

Auto-drafted postmortems


Solves: Saves follow-up work

Approval-gated remediation

Solves: Earned trust

Less searching. More deciding.

ArgusAI gives the on-call engineer a starting point instead of a blank page.



WHAT THIS PROJECT DEMONSTRATES

Problem discovery

User thinking

Product decisions

Metrics

Technical credibility